



Tutorial: SVM In R

袁会卓

Package: e1071

- ▮ Function: `svm()`
- ▮ Usage:
 - ▮ `svmfit<-svm(x,as.factor(y))`
 - ▮ `svmfit<-svm(y ~ ., data=dat, kernel='linear', cost=0.1, scale=FALSE)`
- ▮ Arguments:
 - ▮ `formula, data, x, y, scale, kernel, degree, gamma, cost`

Support vector classifier

- ▮ Generating the observations:
- ▮ `set.seed (1)`
- ▮ `x<-matrix(rnorm(20*2), ncol=2)`
- ▮ `y<-c(rep(-1,10), rep(1,10))`
- ▮ `x[y==1,]=x[y==1,]+1`
- ▮ `plot(x, col=(3-y))`
- ▮ `dat=data.frame(x=x, y=as.factor(y))`

Support vector classifier


- Generating the observations:
- `set.seed (1)`
- `x<-matrix(rnorm(20*2), ncol=2)`
- `y<-c(rep(-1,10), rep(1,10))`
- `x[y==1,]=x[y==1,]+1`
- `plot(x, col=(3-y))`

```
> x
      [,1]      [,2]
[1,] -0.6264538  0.91897737
[2,]  0.1836433  0.78213630
[3,] -0.8356286  0.07456498
[4,]  1.5952808 -1.98935170
[5,]  0.3295078  0.61982575
[6,] -0.8204684 -0.05612874
[7,]  0.4874291 -0.15579551
[8,]  0.7383247 -1.47075238
[9,]  0.5757814 -0.47815006
[10,] -0.3053884  0.41794156
[11,]  2.5117812  2.35867955
[12,]  1.3898432  0.89721227
[13,]  0.3787594  1.38767161
[14,] -1.2146999  0.94619496
[15,]  2.1249309 -0.37705956
[16,]  0.9550664  0.58500544
[17,]  0.9838097  0.60571005
[18,]  1.9438362  0.94068660
[19,]  1.8212212  2.10002537
[20,]  1.5939013  1.76317575
```

```
> y
[1] -1 -1 -1 -1 -1 -1 -1 -1 -1 -1  1  1  1  1  1  1  1  1  1  1
```

Support vector classifier

- Not linearly separable



Support vector classifier

- ▮ Fitting the data:
- ▮ `dat<-data.frame(x=x, y=as.factor(y))`
- ▮ `library('e1071')`
- ▮ `svmfit<-svm(y ~ ., data=dat,`
`kernel='linear', cost=10, scale=FALSE)`
- ▮ `plot(svmfit,dat)`

```
> attributes(svmfit)
$names
[1] "call"           "type"           "kernel"         "cost"
[5] "degree"         "gamma"          "coef0"          "nu"
[9] "epsilon"        "sparse"         "scaled"         "x.scale"
[13] "y.scale"        "nclasses"       "levels"         "tot.nSV"
[17] "nSV"           "labels"         "SV"            "index"
[21] "rho"           "compprob"       "probA"          "probB"
[25] "sigma"         "coefs"          "na.action"      "fitted"
[29] "decision.values" "terms"

$class
[1] "svm.formula" "svm"
```

Support vector classifier

```
> svmfit$index
[1] 1 2 5 7 14 16 17
> summary(svmfit)

Call:
svm(formula = y ~ ., data = dat, kernel = "linear", cost = 10, scale = FALSE)

Parameters:
  SVM-Type:  C-classification
SVM-Kernel:  linear
    cost:    10
   gamma:    0.5

Number of Support Vectors:  7

( 4 3 )

Number of Classes:  2

Levels:
-1 1
```

Support vector classifier


- Cost=0.1:
- `svmfit=svm(y ~ ., data=dat,
kernel='linear', cost=0.1, scale=FALSE)`
- `plot(svmfit,dat)`
- `svmfit$index`
- `summary(svmfit)`

Support vector classifier

cost: 10
gamma: 0.5

Number of Support Vectors: 7


SVM classification plot



cost: 0.1
gamma: 0.5

Number of Support Vectors: 16

SVM classification plot



Support vector classifier

- ▮ Finding optimal tuning parameter:
- ▮ `tune.out <- tune(svm, y ~
., data=dat, kernel="linear", ranges=list(cost=c(0.001, 0.01, 0.1, 1, 5, 10, 100)))`
- ▮ `bestmod = tune.out$best.model`
- ▮ `summary(bestmod)`
- ▮ `bestmod$cost`

```
> attributes(tune.out)
$names
[1] "best.parameters" "best.performance" "method"          "nparcomb"
[5] "train.ind"       "sampling"         "performances"    "best.model"

$class
[1] "tune"
```

Support vector classifier

```
> summary(tune.out)
```

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

cost

0.1

- best performance: 0.1

- Detailed performance results:

cost error dispersion

1 1e-03 0.70 0.4216370

2 1e-02 0.70 0.4216370

3 1e-01 0.10 0.2108185

4 1e+00 0.15 0.2415229

5 5e+00 0.15 0.2415229

6 1e+01 0.15 0.2415229

7 1e+02 0.15 0.2415229

Support vector classifier

- ▮ Construct the test data:
- ▮ `set.seed(1)`
- ▮ `xtest=matrix(rnorm(20*2), ncol=2)`
- ▮ `ytest=sample(c(-1,1), 20, rep=TRUE)`
- ▮ `xtest[ytest==1,]=xtest[ytest==1,] + 1`
- ▮ `testdat=data.frame(x=xtest,
y=as.factor(ytest))`
- ▮ `ypred=predict(bestmod,testdat)`
- ▮ `table(predict=ypred, truth=testdat$y)`

Support vector classifier

- ▮ Construct the test data:
- ▮ `set.seed(1)`
- ▮ `xtest=matrix(rnorm(20*2), ncol=2)`
- ▮ `ytest=sample(c(-1,1), 20, rep=TRUE)`
- ▮ `xtest[ytest==1,]=xtest[ytest==1,] + 1`
- ▮ `testdat=data.frame(x=xtest,
y=as.factor(ytest))`
- ▮ `ypred=predict(bestmod,testdat)`
- ▮ `table(predict=ypred, truth=testdat$y)`

	truth	
predict	-1	1
-1	10	1
1	1	8

Support vector classifier

- ▮ Construct the test data:
- ▮ `set.seed(1)`
- ▮ `xtest=matrix(rnorm(20*2), ncol=2)`
- ▮ `ytest=sample(c(-1,1), 20, rep=TRUE)`
- ▮ `xtest[ytest==1,]=xtest[ytest==`
- ▮ `testdat=data.frame(x=xtest,`
- ▮ `y=as.factor(ytest))`
- ▮ Changing parameter:
- ▮ `svmfit=svm(y~., data=dat, kernel="linear",`
- ▮ `cost=.01, scale=FALSE)`
- ▮ `ypred=predict(svmfit, testdat)`
- ▮ `table(predict=ypred, truth=testdat$y)`

	truth	
predict	-1	1
	-1	10
	1	1
		5

Nonlinear kernel with radial basis


- ▯ Kernels available for svm():
- ▯ linear:
- ▯ Polynomial:
- ▯ radial basis:
- ▯ sigmoid:

Nonlinear kernel with radial basis

```
▮ Construct the test data:  
▮ set.seed (1)  
▮ x=matrix(rnorm(200*2), ncol=2)  
▮ x[1:100,]=x[1:100,]+2  
▮ x[101:150,]=x[101:150,]-2  
▮ y=c(rep(1,150),rep(2,50))  
▮ dat=data.frame(x=x,y=as.factor(y))  
▮ plot(x, col=y)  
▮ train=sample(200,100)  
▮ svmfit=svm(y~., data=dat[train,],  
  kernel="radial", gamma=1, cost =1)  
▮ plot(svmfit , dat[train ,])
```



Nonlinear kernel with radial basis

```
□ Cons  
□ set.  
□ x=ma  
□ x[1:  
□ x[16  
□ y=c(  
□ dat=  
□ plot  
□ trai  
□ svmf  
□ kern  
□ plot
```



Nonlinear kernel with radial basis

```
▮ Constr  
▮ set.se  
▮ x=matr  
▮ x[1:10  
▮ x[10:1  
▮ y=c(re  
▮ dat=da  
▮ plot(x  
▮ train=  
▮ svmfit  
▮ kernel  
▮ plot(s
```



Nonlinear kernel with radial basis

```
> summary(svmfit)
```

Call:

```
svm(formula = y ~ ., data = dat[train, ], kernel = "radial", gamma = 1,  
     cost = 1)
```

Parameters:

```
  SVM-Type:  C-classification  
SVM-Kernel:  radial  
    cost:    1  
   gamma:    1
```

Number of Support Vectors: 37

```
( 17 20 )
```

Number of Classes: 2

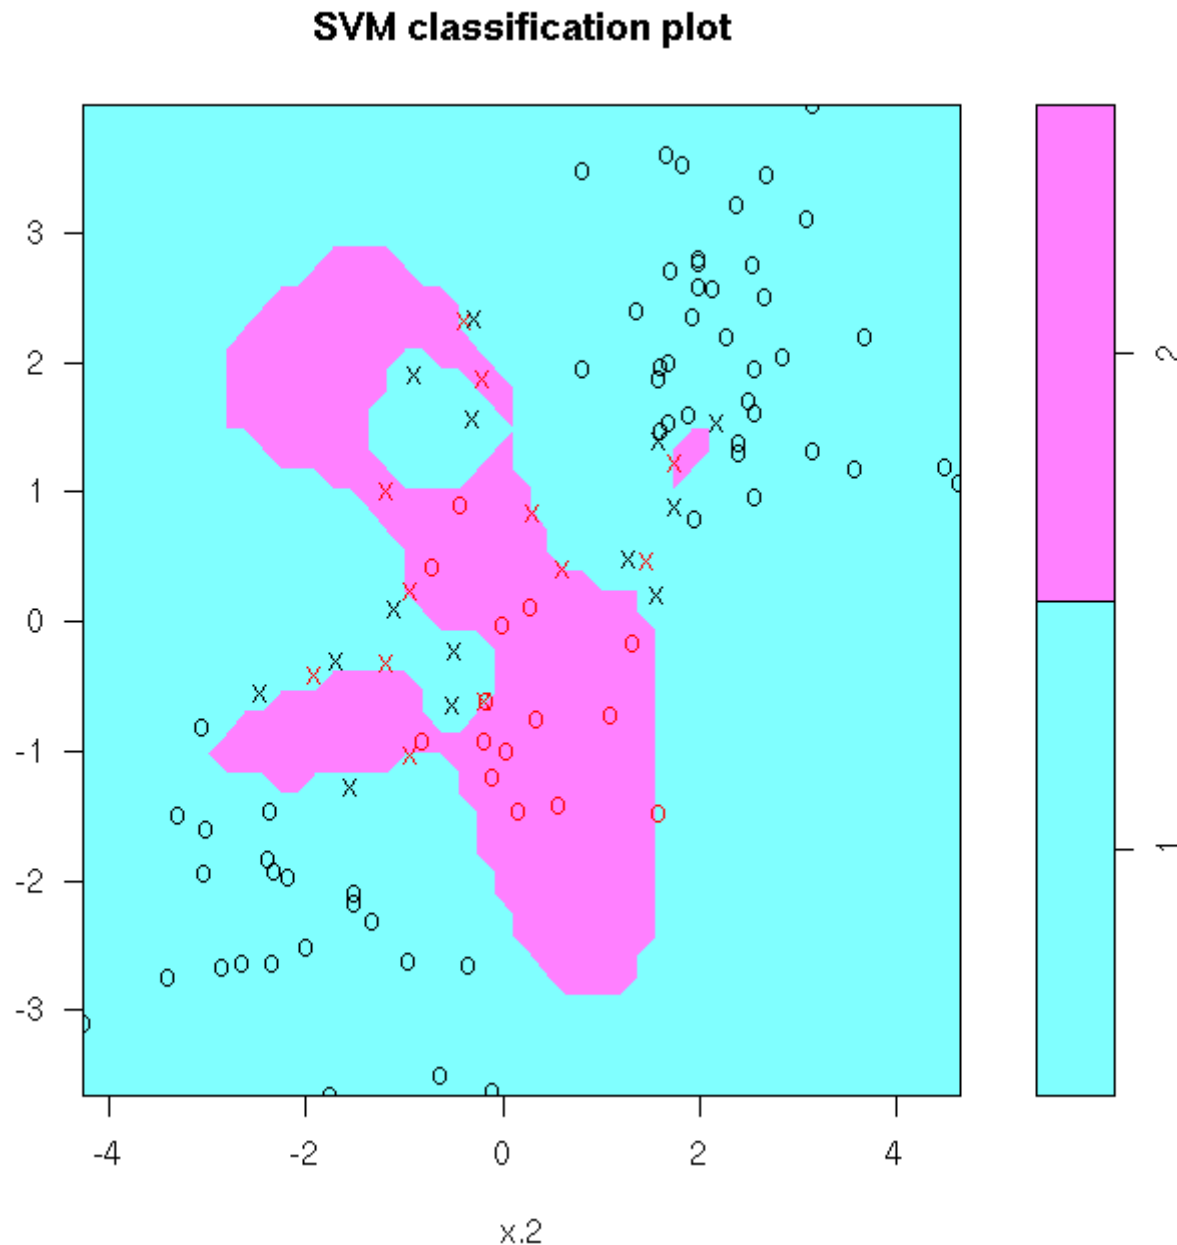
Levels:

```
1 2
```

```
plot(svmfit, dat[train, ],
```

Nonlinear kernel with radial basis

```
Constr  
set.s  
x=mat  
x[1:1  
x[101  
y=c(r  
dat=d  
plot(  
train  
svmfi  
kerne  
plot(  
x.2
```



Nonlinear kernel with radial basis

- ▮ Choosing parameter, cost=1, gamma=0.5:
- ▮ `tune.out=tune(svm, y~., data=dat[train,],
kernel="radial",
ranges=list(cost=c(0.1,1,10,100,1000),
gamma=c(0.5,1,2,3,4)))`
- ▮ `table(true=dat[-train,"y"],
pred=predict(tune.out$best.model,
newx=dat[-train ,]))`

```
      pred  
true   1   2  
  1  54  23  
  2  17   6
```

Nonlinear kernel with radial basis

```
> summary(tune.out)

Parameter tuning of 'svm':

  sampling method: 10-fold cross validation

  best parameters:
    cost gamma
      1    0.5

  best performance: 0.13

Detailed performance results:
  cost gamma error dispersion
1  1e-01   0.5  0.28 0.16865481
2  1e+00   0.5  0.13 0.10593499
3  1e+01   0.5  0.14 0.11737878
4  1e+02   0.5  0.18 0.10327956
5  1e+03   0.5  0.19 0.11005049
6  1e-01   1.0  0.30 0.15634719
7  1e+00   1.0  0.14 0.11737878
8  1e+01   1.0  0.18 0.10327956
9  1e+02   1.0  0.19 0.13703203
```

ROC curves

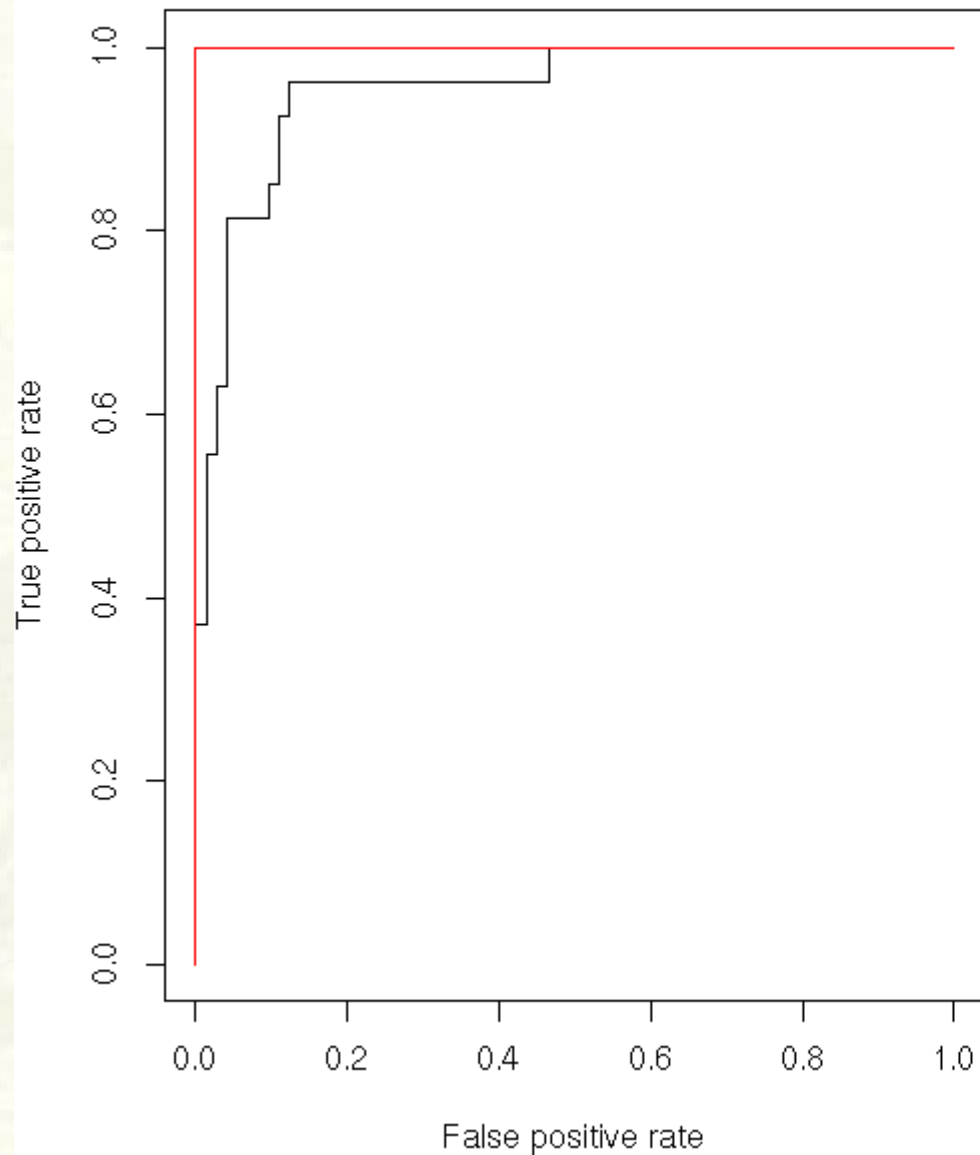
- ▮ Define function to plot ROC curves:
- ▮ `rocplot=function(pred, truth, ...){`
- ▮ `predob = prediction(pred, truth)`
- ▮ `perf = performance(predob, "tpr", "fpr")`
- ▮ `plot(perf,...)}`
- ▮ `pred: numerical score`
- ▮ `truth: true label`

ROC curves


- ▮ Plot ROC curves:
- ▮ `svmfit.opt=svm(y~., data=dat[train,],
kernel="radial", gamma=2,
cost=1,decision.values=T)`
- ▮ `fitted=attributes(predict(svmfit.opt,dat[t
rain,],decision.values=TRUE))
$decision.values`
- ▮ `par(mfrow=c(1,2))`
- ▮ `rocplot(fitted,dat[train,"y"],main="Traini
ng Data")`

ROC curves

Training Data



Test Data




SVM with multiple classes


```
▮ SVM with multiple classes:  
▮ set.seed (1)  
▮ x=rbind(x, matrix(rnorm(50*2), ncol=2))  
▮ y=c(y, rep(0,50))  
▮ x[y==0,2]=x[y==0,2]+2  
▮ dat=data.frame(x=x, y=as.factor(y))  
▮ par(mfrow=c(1,1))  
▮ plot(x,col=(y+1))  
▮ svmfit=svm(y~., data=dat, kernel="radial",  
  cost=10, gamma=1)  
▮ plot(svmfit, dat)
```

SVM with multiple classes

```
library(SVM)
set.seed(123)
x=rbinom(100,1)
y=c(y=rbinom(100,1))
x[y==0]=rbinom(50,1)
dat=data.frame(x,y)
par(mfrow=c(2,2))
plot(x,y)
svmfit=svm(x,y)
plot(x,y)
```

SVM classification plot





Thank you!